

Adversarial Text Generation

NLP and Deep Learning — Final Project
KSNLPDL2KU

Andreas Holck Høeg-Petersen
anhh@itu.dk

Mathias Bastholm
mbas@itu.dk

August 14, 2020

1 Introduction

In recent years, Generative Adversarial Networks (GANs) have gained a lot of traction in the Deep Learning community because of their impressive results in image generation. The general idea is that a generator and a discriminator are jointly trained to produce an image output that is seemingly indistinguishable from non-generated images. This model were first described in Goodfellow et al. 2014.

We want to attempt to apply this strategy for text generation. The main difficulty for this task is that whereas image outputs can be considered a continuous value, a sentence is inherently discrete as it is a sequence of words each of which is chosen by the model using the non-differentiable *argmax* function. To remedy this, we propose a model where the discriminator is trained to distinguish between the continuous outputs of a pre-trained encoder given a ‘true’ sentence from the generated, ‘fake’ output stemming from our generator.

In our project, we will construct and train an autoencoder model that can encode and decode a sentence from English to English. The encoded sentences are then used as labelled training data for the discriminator, representing ‘true’ values. The job of the generator is to produce similar encodings but doing this from random noise in a way that makes the discriminator unable to distinguish between the encodings stemming from the autoencoder and the encodings stemming from the generator.

Ideally, this would train the generator to produce sentence encodings that can be fed to the decoder of the Transformer model which would then produce meaningful sentences from this artificially generated input. See Figure 1 for an overview of the complete model.

This project thus have two objectives: one is to construct a working autoencoder that can map an English sentence to some hidden state \mathbf{X} with a corresponding decoder that can extract the original sentence from \mathbf{X} . For convenience, we will refer to the encoder part of this model as the ‘Teacher’. The second objective is to build a GAN network, where a generator — the ‘Student’ — must learn to produce approximations of \mathbf{X} .

The second objective is highly experimental as explained in Section 2, where we will also describe other approaches at using the GAN architecture for NLP problems. We will then describe how we have build the different parts of the model and how we utilize our dataset. Then we will present our results and discuss the shortcomings of the models, and proceed to suggest improvements and ideas for further research. Lastly we provide a conclusion on the project.

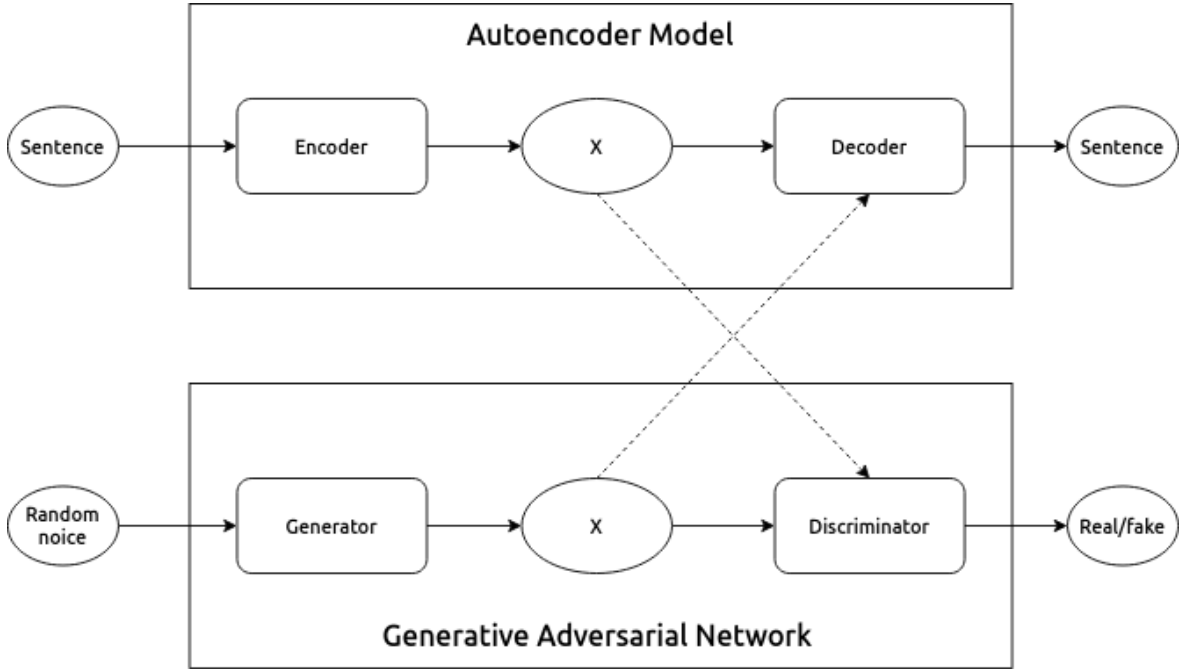


Figure 1: Overview of the model architecture. The dotted lines from the \mathbf{X} s represents that the encoded and generated \mathbf{X} s will be fed to the discriminator and the decoder during training and evaluation, respectively.

2 Background

Applications GANs have mainly focused on image generation and has not yet seen a major breakthrough in text generation. As mentioned in Section 1, this is because the discrete nature of text, which is basically a sequence of words, requires a non-differentiable *argmax* function to transform a probability distribution over a vocabulary to a single value (ie. the word with the highest probability). This is depicted in Figure 2.

There have been, however, multiple attempts at working around this issue. According to Chintapalli 2019 these approaches can broadly be categorized into three types:

- Reinforcement Learning-based solutions
- The Gumbel-Softmax approximations
- Avoiding discrete spaces by working with the continuous output of the generaor

In this project, we follow the third approach, but in this section we will give short introductions to the idea behind the two first.

As an example of an RL-based solution, Yu et al. 2016 proposes the SeqGAN model, in which the generator is considered an RL-agent with states \mathbf{s}_t being the text generated at timestep t and actions \mathbf{a} being all the possible words to choose next. The agent then chooses its next word (takes an action a) based on some policy function $\pi(a | \mathbf{s}_t, \theta)$, where θ are the parameters to be optimized. Using Monte-Carlo rollouts to produce a number of different sentences sharing a prefix \mathbf{s}_t , the discriminator then rewards each sentence and the averaged reward is then used to perform gradient ascent on $\mathbf{J}(\theta)$, where \mathbf{J} is the performance measure of θ . This approach alleviates some of the problems of training

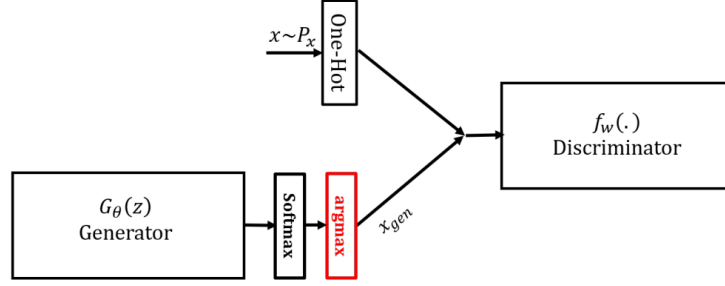


Figure 2: A simple GAN model where the generator output is run through an *argmax* function before being given to the discriminator. This prevents gradients to flow from the discriminator to the generator. Source: Haidar and Rezagholizadeh 2019.

a GAN for text generation, but it suffers from an unstable and slow training process, convergence to sub-optimal local minima and and extremely large state-space.

Another approach is to use the Gumbel-Softmax distribution to approximate a one-hot encoding of a probability distribution passed through the *argmax* function. This is the approach taken by Kusner and Hernández-Lobato 2016. Here, a d -dimensional one-hot encoding vector \mathbf{y} is approximated using

$$\mathbf{y} = \text{soft max}\left(\frac{1}{\tau}(\mathbf{h} + \mathbf{g})\right) \quad (1)$$

where \mathbf{h} is some hidden state (ie. of an RNN), \mathbf{g} is drawn from a Gumbel distribution and τ is a temperature parameter. This works because it is differentiable and as $\tau \rightarrow 0$ the distribution of \mathbf{y} will match that we get from

$$\mathbf{y} = \text{one_hot}\left(\arg \max_i (h_i + g_i)\right) \quad (2)$$

which again can be shown to be the same as sample \mathbf{y} from a probability distribution $\mathbf{p} = \text{softmax}(\mathbf{h})$ where $p_i = p(y_i = 1), i = 1 \dots d$.

Finally, for other examples on the third approach, see Donahue and Rumshisky 2018 and Haidar and Rezagholizadeh 2019.

3 Method

For all our models (autoencoders and GANs), we drew inspiration PyTorch tutorials (Robertson 2020, *Sequence-to-Sequence Modeling with nn.Transformer and TorchText* 2020, Inkawhich 2020), tweaking them to our specific needs. The whole process (project development, research, data collection, coding, training, experimentation and analysis) was conducted during a 10-day period.

This section will describe this process, focusing on the final outcomes rather than including all our intermediate steps and missteps.

3.1 Dataset

For training of the autoencoders, we simply needed dataset consisting of a large number of English sentences. We obtained this from Universal Dependencies, where used the ‘Universal Dependencies — English Dependency Treebank Universal Dependencies English Web Treebank v2.6 — 2020-05-15’ (Silveira et al. 2014) consisting of 12,543 training sentences, 2,077 test sentences, 2,002 dev sentences and a vocabulary of 16,654 training tokens. The data is annotated with metadata such as lemmas and word classes, but we discarded this information as it was not relevant for our purpose.

Furthermore, we also utilized another dataset intended for training English-to-French translation. This dataset originates from <https://tatoeba.org/eng/> and consists of 135,842 sentence pairs. We discarded the French sentences and removed all duplicate sentences and sentences of length smaller than 3, as well as splitting all sentences that contained punctuations, questionmarks, exclamation points, etc. This gave us a set of 92,343 sentences, which we split 80/10/10 between training and training/validation/testing. The dataset had a vocabulary of 13,731 tokens.

3.2 Models

We developed two different versions of the autoencoder model and one GAN model. These will be described in this subsection.

3.2.1 The TransformerModel

Our first autoencoder was based on *Sequence-to-Sequence Modeling with nn.Transformer and Torch-Text* 2020. This model consists of a very simple decoder, that is simply a feed-forward neural net that takes a 2-dimensional tensor $X \in \mathbb{R}^{n \times k}$ and maps each of the n k -dimensional vectors of the sequence to a probability distribution over the entire vocabulary which it then can convert to an output sequence using *argmax*.

The encoder, however, is responsible for generating X and it does so by using the Transformer architecture as suggested in Vaswani et al. 2017 and implemented in the PyTorch module `nn.Transformer`. For our purpose, however, we only used the submodule `nn.TransformerEncoder`, which consists of a stack of encoder layers that uses self-attention to focus on specific, relevant parts of the input sequence in one go, and then passes its output on to the next layer through a feed-forward network. As a preprocessing step, before the input sequence is passed through the transformer, positional encoding is added, as suggested in the paper (ibid.).

This model also uses an embedding as its first layer. For embeddings, we use pretrained word vectors from the `polyglot` Python package. These have an embedding size of 64, which therefore what we use across all models. Note that the next, RNN-based autoencoder uses the same embedding setup.

3.2.2 RNN and Attention based Autoencoder

In our second model, the heavy-lifting is switched from the encoder to the decoder. Again, we utilize the attention mechanism, but this time it is combined with an RNN architecture, more precisely a Gated Recurrent Unit (GRU). In this model, the encoder is simply a GRU layer that processes the entire sequence and then its final hidden state aswell as the output for each word in the sequence is passed to the decoder.

The decoder has a bit more to it. On each timestep it takes in its own last output (starting with the special **start-of-sequence-token**), a hidden state (starting with the last hidden state of the encoder)

and all the encoder outputs. It then uses a linear layer (ie. a simple feed-forward network) to calculate the attention weights by combining the input and the current hidden state. The attention weights and the encoder outputs are then multiplied together using matrix multiplication and the result of this operation can then be merged with the original input and passed through another linear layer to produce a vector of size `(sequence_length * hidden_size)`.

In this vector, the decoder has now embedded all the information about where to focus its attention, and it can pass this to a GRU just as in the encoder — however, opposite to the GRU in the encoder, the output of the recurrent layer in the decoder is responsible for mapping back to actual words. This mapping is finalized by an output linear layer that expands the hidden size to the size of the vocabulary and then, finally, a *softmax* operation converts the output to probability distributions.

When the decoder outputs the special **end-of-sequence**-token it is finished and the process terminates. Our implementation of this model is based on Robertson 2020.

3.2.3 GAN

Our GAN model is modelled after our TransformerModel in an attempt to mimic the inner mechanics of the network it is trying to imitate. As stated, a GAN consists of a generator and a discriminator. The generator gets a vector of random noise as input. This vector has dimensions `(max_sequence_length * embedding_size)`. The reason we set a parameter `max_sequence_length` is because neural networks expects static sizes, but since sentences can have different lengths, we set an upper bound. The generator should, however, be able to encode its output in a way, that produces sentences of length 0 to `max_sequence_length` (ie. by encoding an **end-of-sequence**-token at some position).

The generator passes its input through a linear layer and then, as the TransformerModel, adds positional encoding before it runs it through an encoder stack (`nn.TransformerEncoder`). The final output has the same shape as the output of the encoder part of the TransformerModel.

The discriminator has almost the same structure as the generator (it uses positional encoding and an encoder stack), but the input differs and it outputs a single number between 0 and 1. Recall, that the job of the discriminator is to decide whether its input was generated from random noise via the generator or if it was an encoding stemming from a real sentence passed through a trained encoder. Thus, it is effectively a binary classifier, and its output should describe its conviction that the input is 'true'.

3.3 Training

Our training had two aspects: first, we had to train our autoencoders on English-to-English sentence pairs, and then we had to train our GAN model. This section will describe both.

3.3.1 Training the autoencoders

This part is pretty straight forward. We give the model a batch of sentences which it then encodes and decodes again. Since the input sentence is also the target sentence, the loss is simply the difference between the input and the output. To calculate this loss, we use PyTorch's `nn.CrossEntropyLoss` which takes the raw output of the model (ie. before it has been converted to a probability distribution) and a target vector and then calculates the negative log-likelihood loss. We use Stochastic Gradient Descent to update the weights of the models.

There were some notable differences between the way we trained the TransformerModel and the RNN-based model. For the TransformerModel we had good results with a learning rate of 0.05, and we performed training on the entire training set in batches of 8 and over the course of 25 epochs. We also trained it another time with just 5 epochs, but this time without ignoring the padding in the batches (when batching the data, not all sentences have the same length and the smaller sentences are padded with the special <PAD>-token). This makes the model achieve a lower loss much faster, as it quickly learns the pattern of the padding tokens, but this often comes at the expense of the model learning the actual sentence, which generally is what is of interest. However, we wanted our Student model to learn both from a Teacher that had learned padding and one that hadn't, so we trained the autoencoders (Teacher) on both.

In training the RNN-based model we took inspiration from Robertson 2020. This means we set the learning rate to 0.01 and instead of training on the entire training set over multiple epochs, we let the model randomly pick a data batch, train on that and then pick a new for 60,000 iterations. Furthermore, we used teacher forcing, where instead of letting the output of the decoder be its own next input, we give the actual next input (from the target sentence). This approach can lead to some terrible cases of overfitting and preventing the model from learning its mistakes properly, but in the right dose it can also help the model converge faster. Therefore, for each training instance, we chose with 50% probability whether to use teacher forcing or not.

3.3.2 Training the GAN model

The training of the GAN model is the interesting and difficult process. Here, we have to simultaneously train two different models that work towards opposite goals. One, the generator, will have to learn how to maximise the probability that the discriminator classifies its output as 'true', while the other, the discriminator, will have to learn to maximise the probability that it classifies the output of the generator as 'fake' while also classifying the output of the Teacher (encoder) as 'true'. So the two components engage in a min-max game with one another.

Formally, let G and D be the generator and the discriminator respectively and let x be the data representing a sentence (encoded or generated). Then, $D(x)$ is the output of the discriminator given some data, and we want this to be high when x is 'true' (stemming from the Teacher) and low when x is 'fake' (generated by the Student). Let z be some random noise, and then we have that $G(z)$ is the output generated by the Student given some random noise. So, the discriminator now tries to *maximize* the probability that it correctly classifies 'true' and 'fake' data, $\log D(x)$, and the generator tries to *minimize* the probability that the discriminator classifies its outputs as 'fake', $\log(1 - D(G(z)))$. As described in the original GAN paper Goodfellow et al. 2014, this can be formalised as:

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(x)} [\log(1 - D(G(z)))]$$

where p_{data} is the 'true' distribution over the real data which the generator tries to estimate, and p_z is the estimated distribution that the generator draws its samples from.

The way we do this is by following the approach suggested by *ibid.* and PyTorch GAN tutorial for image generation Inkawhich 2020. This involves first giving the discriminator a batch of 'real' input (ie. input generated by the Teacher) and then calculating the loss and the gradients based on its performance. Then we give it a batch of 'fake' data generated by the Student, again calculating the loss and gradients. Then, before we move on to training the generator, we perform our optimization step for the discriminator. When we then pass the output of the generator through the discriminator, the discriminator has been updated and the generator will have to learn to adjust to this.

However, one could imagine that the generator could learn to just output random words very well and that the discriminator would not be able to distinguish this from a ‘real’ sentence. Therefore, before training the generator, we also give the discriminator a batch of ‘fake’ data that has been generated by the encoder, but when the encoder have gotten random words as input. This, in theory, should force the generator to produce meaningful sentences (not just random words) for it to be classified as ‘true’ by the discriminator. This adjustment is specific to our approach.

It is worth noting, that training a GAN is known to be highly non-trivial (Hui 2018) and even small variations in the hyperparameters, number of training iterations or internal structure of the generator and/or discriminator can make the model collapse spectacularly.

4 Results and analysis

In this section, we will present and discuss our results as well as give suggestions for further improvements and research.

4.1 TransformerModel

As stated, we trained the TransformerModel both with and without ignoring padding. Generally, we achieved nice results here, with a word accuracy around 99%. In Figure 3 and 4 we have plotted the loss.




Figure 3: Loss of the standard TransformerModel




Figure 4: Loss of the TransformerModel that takes padding into account and has EOS tokens. Even with way fewer epochs, the results were very good.

The following are some examples of how the model performed on specific inputs:

```
++++++ Successful encoding ++++++
> he can play the piano , the flute , the guitar , and so on .
< he can play the piano , the flute , the guitar , and so on .

> i have to take the entrance examination today .
< i have to take the entrance examination today .

++++++ Failed encoding ++++++
```

```
> i'm sure it wouldn't be too hard to find out who hacked into our system .
< i'm sure it wouldn't be too hard to find out who weakly into our system .

> you have the right to free speech , but not the right to slander .
< you have the right to free speech , but not the right to limits .
```

And here are some examples of the model when train with padding and **end-of-sequence-token**:

```
> his house was small and old . <EOS>
< his house was small and old . <EOS>

> i realized what was happening . <EOS> <PAD>
< i realized what was happening . <EOS> <PAD>

> there is no museum . <EOS> <PAD> <PAD> <PAD>
< there is no antidote . <EOS> <PAD> <PAD> <PAD>
```

Notice the error in the last sentence. For good measure, we report the following input/output-pair that was observed during training. Maybe the model were developing a political opinion, maybe it was joking. We will probably never know:

```
> rwandan rebels are pushing their offensive south as fighting continues in the capital kigali . <EOS>
< honest guys are noticed their interesting south as fighting awake in the capital fortune . <EOS>
```

4.2 RNNModel

Our RNN-based model performed much worse, than we had expected. It also took an extraordinary amount of time to train, so we only did so many experiments. A plot of the loss can be seen in Figure 5.

We trained the same model several times with decreasing batch size, and even though the loss kept going down, the following examples show that the end result was not amazing:

```
> tom got to the station too late so he missed the train . <EOS> <PAD>
< tom got to the so so so he opened the same . <EOS>

> i don't want to waste my time trying to do this again . <EOS> <PAD>
< i don't want to take my job to do me it . <EOS>

> if the car is gone , he can't be at the office . <EOS> <PAD>
< if the book is he , he was late than the . <EOS>

> i can't believe that you were the smartest kid in your class . <EOS> <PAD>
< i can't believe that you were a a the in the . . <EOS>
```

It has clearly learned something, most notably **end-of-sequence-tokens** and padding. Further, it seems that it does a fine job in the beginning of sentences, but then loses track of what to do. This indicates that it hasn't picked up on long-term dependencies in the sentences, which might could be alleviated with another RNN-unit (eg. an LSTM) or more training.




Figure 5: Training loss of RNN-based attention model.

4.3 GAN

Now, let's look at the results of our GAN model. Because of the poor performance of our RNN-based autoencoder and some technical difficulties, we only trained the GAN against the TransformerModel. Also, we set the maximum sequence length to be 10 which is fine for prototyping and evaluating if anything is actually learned.

Beginning with Figure 6 and Figure 7 we see how the network performed when trained with Adam and Stochastic Gradient Descent respectively as its training algorithm. The difference is very clear: for Adam, the generator and the discriminator has a pretty consistent relationship until suddenly, the discriminator seems to figure out how to call bluff on the generator, and the generator has no answer to that. The loss goes way up for the generator and it falls to a minimum for the discriminator and the model seems to collapse.

For SGD, the model is in some sense way more stable and the loss for both the discriminator and the generator stays in same limited interval.




Figure 6: GAN loss using Adam




Figure 7: GAN loss using SGD

Since what we really wanted was a strong generator, and this seemed to be the component that failed, we tried training with a double batch for the generator to give an edge against the discriminator. The resulting loss is seen in Figure 8 and Figure 9. The pattern is the same, but this time, the Adam model collapses faster.




Figure 8: GAN loss using Adam and double batch for generator




Figure 9: GAN loss using SGD and double batch for generator

Looking at the actual output, we some interesting results. First, lets look at some output from the Adam training (without double batch):

```
[0/10] [0/6765]      Loss_D: 1.8526  Loss_G: 0.6623  D(x): 0.5776    D(G(z)): 0.5878 / 0.5499
know for know i'm i'm the for good some i'm
```

```
[0/10] [500/6765]    Loss_D: 1.7199  Loss_G: 0.6846  D(x): 0.2188    D(G(z)): 0.5105 / 0.5054
your appropriated make ? but <PAD> <EOS> <EOS> <PAD> <PAD>
```

```
[0/10] [2000/6765]   Loss_D: 1.2057  Loss_G: 0.9821  D(x): 0.2932    D(G(z)): 0.4748 / 0.3959
your out out out out out <EOS> <EOS> <PAD> <PAD>
```

...

[2/10] [2450/6765] Loss_D: 1.4581 Loss_G: 1.7943 D(x): 0.2807 D(G(z)): 0.1915 / 0.1928
 tom cats a a stayed . <EOS> <PAD> <PAD> <PAD>

[2/10] [2950/6765] Loss_D: 0.4975 Loss_G: 1.7132 D(x): 0.3910 D(G(z)): 0.0857 / 0.2305
 he who painted is is . <EOS> <PAD> <PAD> <PAD>

[2/10] [3450/6765] Loss_D: 0.4640 Loss_G: 6.3138 D(x): 0.3933 D(G(z)): 0.1249 / 0.0479
 don't child caught more my . . <EOS> <PAD> <PAD>

These are outputs that show, what the loss of the discriminator (Loss_D), the loss of the generator (Loss_G), the average discriminator output on non-generated data ($D(x)$) and the average discriminator output on the generated data before and after updating the discriminator ($D(G(z))$).

When we train the discriminator on encoded random words, $D(x)$ should be 0.5 for a perfect D (because half of x is 'true' and the other half is 'fake'). When omitting the encoded random words, $D(x)$ should be 1. Of course, for a perfect network with a perfect D and a perfect G , both $D(x)$ and $D(G(z))$ should be 0.5.

We see that the model learns the structure of **end-of-sequence**-tokens, padding and to some extent punctuation. The sentences are clearly nonsense, but they do not seem totally random. Compare the later outputs with the first and these points are more clear.

We also notice, that $D(G(z))$ become very low towards the end here. This means that the discriminator correctly classifies nearly all the generated sentences.

Now let's look at the output for the SGD training:

[0/10] [0/6765] Loss_D: 1.6203 Loss_G: 1.6716 D(x): 0.2540 D(G(z)): 0.2270 / 0.2125
 . where ? where long where where say long where

[0/10] [500/6765] Loss_D: 1.3528 Loss_G: 1.0269 D(x): 0.3616 D(G(z)): 0.4392 / 0.3740
 <EOS> <EOS> <EOS> do is <EOS> <EOS> <EOS> <PAD> <PAD>

[0/10] [3000/6765] Loss_D: 1.4074 Loss_G: 0.8827 D(x): 0.2683 D(G(z)): 0.3572 / 0.4423
 . why the this this <PAD> <PAD> <PAD> <PAD> <PAD>

...

[1/10] [6200/6765] Loss_D: 1.3663 Loss_G: 0.9820 D(x): 0.2959 D(G(z)): 0.4286 / 0.4028
 . a he a <EOS> <EOS> <PAD> <PAD> <PAD> <PAD>

[1/10] [6700/6765] Loss_D: 1.0762 Loss_G: 0.9794 D(x): 0.3373 D(G(z)): 0.4294 / 0.4404
 . a he why <EOS> <EOS> <EOS> <PAD> <PAD> <PAD>

[2/10] [450/6765] Loss_D: 0.8063 Loss_G: 0.8995 D(x): 0.3487 D(G(z)): 0.3294 / 0.4373
 . a he a <EOS> <EOS> <EOS> <PAD> <PAD> <PAD>

Though we don't see the collapse in the generator, the output sentences are much worse from a human perspective. It does put padding and **end-of-sequence**-tokens in their right place, but it uses too many and it doesn't seem to get punctuation. Also, it converges towards a smaller and smaller sentence with a still more limited vocabulary.

Next up, we have some examples of output when training with Adam and double batch for the generator. This time, the first output sentence is based on some fixed noise initialised before training, and the second sentence comes from random noise generated at each print:

[0/10][0/6765] Loss_D: 1.6505 Loss_G: 0.5963 D(x): 0.5498 D(G(z)): 0.5673 / 0.5897
. . . . <EOS> . . . common likes
became <EOS> <EOS> <EOS> . <EOS>

[0/10][500/6765] Loss_D: 1.6686 Loss_G: 0.6725 D(x): 0.2467 D(G(z)): 0.5252 / 0.5183
? ? <PAD> <PAD> about something <EOS> <PAD> <PAD> <PAD>
? ? <PAD> of about martial make <PAD> <PAD> <PAD>

[0/10][3000/6765] Loss_D: 1.4506 Loss_G: 0.7180 D(x): 0.2317 D(G(z)): 0.4720 / 0.4948
? ? apartment thing he <EOS> <EOS> <EOS> <PAD> <PAD>
? my apartment thing but <EOS> <EOS> <PAD> <PAD> <PAD>

...

[2/10][6450/6765] Loss_D: 0.9028 Loss_G: 1.7122 D(x): 0.3388 D(G(z)): 0.1354 / 0.1937
into i've me him allow . <EOS> <PAD> <PAD> <PAD>
into i've necessary him allow . <EOS> <PAD> <PAD> <PAD>

[3/10][200/6765] Loss_D: 0.5534 Loss_G: 4.4299 D(x): 0.4601 D(G(z)): 0.1967 / 0.0788
days days still going going . . <EOS> <EOS> <PAD>
days was still my going . . <EOS> <EOS> <EOS>

[3/10][700/6765] Loss_D: 0.7864 Loss_G: 1.9832 D(x): 0.3432 D(G(z)): 0.0988 / 0.1549
smoking his get his need careful . <EOS> <PAD> <PAD>
smoking a get get probably careful . <EOS> <PAD> <PAD>

Here we also observe a generator that looses to the discriminator, but the output sentences are not that bad. Interestingly though, both the fixed and the random noise produces almost similar results after the first output which may indicate that the model tries to find *one* correct sentence to give to the discriminator.

For good measure, we also tried to see what happened if we didn't train the discriminator on random encoded words (as described in Section 3.3.2). As can be seen in Figure 10 this let to a model that for long was much more stable than the previous ones and with a generator having a lower loss than the discriminator. However, at a certain point, the model basically explodes and the training completely breaks.

Here are some outputs of that model:

[0/10][0/6765] Loss_D: 1.9186 Loss_G: 0.6208 D(x): 0.4685 D(G(z)): 0.5844 / 0.5884
shorter anything anything out go again out mary out prefer
machine shopping fish it's shopping prefer mistakes mistakes everything clear

[0/10][500/6765] Loss_D: 1.4489 Loss_G: 0.7317 D(x): 0.5092 D(G(z)): 0.5200 / 0.4889
. ? of <PAD> <PAD> of of way she <PAD>
. ? of <PAD> <PAD> of of english she <PAD>

[0/10][1000/6765] Loss_D: 1.5700 Loss_G: 0.6301 D(x): 0.4740 D(G(z)): 0.5423 / 0.5395
<PAD> untidy of <PAD> <PAD> <PAD> of way she <PAD>
<PAD> make of <PAD> <PAD> poor of way his <PAD>

...




Figure 10: Loss of the GAN with double batching for the generator and no encoded random words for the discriminator.

```
[1/10][5200/6765]    Loss_D: 1.4087    Loss_G: 0.7026    D(x): 0.5035    D(G(z)): 0.5127 / 0.4960
<PAD> appropriated she she she <EOS> <EOS> <EOS> <PAD> <PAD>
<PAD> appropriated she she she <EOS> <EOS> <EOS> <EOS> <PAD>
```

```
[1/10][5700/6765]    Loss_D: 1.4143    Loss_G: 0.7011    D(x): 0.4879    D(G(z)): 0.5002 / 0.4968
worms <EOS> <PAD> ? why <PAD> <EOS> <EOS> <PAD> <PAD>
worms <EOS> worms <PAD> something <PAD> <PAD> <EOS> <PAD> <PAD>
```

```
[1/10][6200/6765]    Loss_D: 1.3873    Loss_G: 0.6957    D(x): 0.5044    D(G(z)): 0.5023 / 0.4997
? <EOS> do <PAD> you'll thing <EOS> <EOS> <PAD> <PAD>
? <EOS> do <PAD> of thing <EOS> <EOS> <EOS> <PAD>
```

Here we seem to see that model has much harder time constructing actual sentences. It puts **end-of-sequence-tokens** and padding at random places and it hasn't really picked up on punctuation yet. But we are not yet at the collapse seen in Figure 10, so let's take a look at some more outputs:

```
[2/10][1950/6765]    Loss_D: 1.3705    Loss_G: 0.7144    D(x): 0.5072    D(G(z)): 0.4975 / 0.4904
get appropriated of is <EOS> <EOS> <EOS> delayed <PAD> <PAD>
get appropriated of is <EOS> <EOS> <EOS> a <PAD> <PAD>
```

```
[2/10][2450/6765]    Loss_D: 1.3432    Loss_G: 0.6959    D(x): 0.5086    D(G(z)): 0.4853 / 0.4991
that criticisms of to that <EOS> <EOS> <EOS> <PAD> <PAD>
```

that criticisms of him <EOS> <EOS> <EOS> <EOS> <EOS> <PAD>

[2/10][2950/6765] Loss_D: 1.4055 Loss_G: 0.6989 D(x): 0.4902 D(G(z)): 0.4984 / 0.4975
is much not to <EOS> <PAD> <PAD> <PAD> <PAD> <PAD>
is my not <EOS> <EOS> <PAD> <PAD> <PAD> <PAD> <PAD>

...

[3/10][200/6765] Loss_D: 1.3366 Loss_G: 0.7089 D(x): 0.5265 D(G(z)): 0.4985 / 0.4929
. . . . <EOS> <EOS> <EOS> <PAD> <PAD> <PAD>
. . . . <EOS> <EOS> <EOS> <PAD> <PAD> <PAD>

[3/10][700/6765] Loss_D: 0.0100 Loss_G: 5.0729 D(x): 0.9973 D(G(z)): 0.0073 / 0.0072
<PAD> <PAD> <PAD> <PAD> <PAD> <PAD> <PAD> <PAD> <PAD> <PAD>
<PAD> <PAD> <PAD> <PAD> <PAD> <PAD> <PAD> <PAD> <PAD> <PAD>

[3/10][2700/6765] Loss_D: 0.0017 Loss_G: 6.7365 D(x): 0.9996 D(G(z)): 0.0013 / 0.0013
make make make make make make make make make make
make make make make make make make make make make

Here it seems that the model was beginning to pick up on where to put **end-of-sequence-tokens** and padding (though not to what extent) but when the collapse happens, the generator completely stops working and just outputs single-word sentences. The collapse is also clearly represented in both $D(x)$ and $D(G(z))$.

Finally, for comparison, lets look at some examples of the GAN model, when it is trained on a TransformerModel that hasn't learned about padding and **end-of-sequence-tokens**:

are thank new to you . suspicious suspicious suspicious prevents suspicious
mafia mafia mafia subtitled

something troll me watched watched . suspicious suspicious suspicious suspicious
mafia mafia mafia mafia mafia

your ignored delayed delayed . helpless helpless helpless helpless helpless
helpless suspicious helpless helpless helpless

it's , , him looking expect . . housewarming housewarming housewarming
suspicious mercury mercury mercury

for change but do i . suspicious suspicious suspicious suspicious suspicious
suspicious suspicious mafia mafia

i mohammed convenience play the . suspicious suspicious suspicious suspicious
suspicious mafia mafia mafia mafia

he used was donated money . subtitled suspicious suspicious suspicious
suspicious suspicious suspicious scientist scientist

It is interesting note, that it seems to attempt generating a sentence with different word and some kinds of sentence constructs, but after it outputs a punctuation mark it almost just outputs the same word over and over. This indicates that it has learned the concept of sentence length and some

consequence of an end of sequence marker (such as the punctuation mark).

4.4 Further research and improvements

We believe that the full potential of using the TransformerModel encodings in a GAN has yet to be revealed. Training a GAN without batching using a variable input length with EOS tokens, could lead to further improvements. Another approach to explore is using the original transformer encoder-decoder stack in place of the two encoder stacks approach we used.

Furthermore using a fixed size encoding such as the RnnModel makes training the encoder much harder but should make training the GAN much easier. As such getting a good Teacher model using this approach could potentially lead to large improvements in the GAN model. To this, it is also worth considering, that where some encoders might just learn to map words to some hidden state, a better encoder could probably inject syntax and sentence-meaning, which might also help a GAN model perform better.

There might also be advantages in systematically testing how training evolves with different optimizers, such as Adam, SGD, Adagrad, etc. and with different learning rates. The GAN model is clearly very fragile and it seems that very careful tuning of hyperparameters is needed to avoid a collapse. Maybe it would even be beneficial with dynamic hyperparameters that adapts to the state of the GAN.

Further, it would also be beneficial to get a better metric for evaluating the performance of the GAN. The loss and the output of the discriminator gives a hint to how the network is doing, but often it does not directly relate to how seemingly ‘good’ sentences produced by the generator of the model.

5 Conclusion

We have explored two different ways of encoding sentences — one, using a small Transformer encoder stack and a simple feed-forward decoding network, and another using an RNN-based encoder and a more involved attention-based decoder. We achieved good results with the first, but the second one took extremely long to converge and only got to sub-par results.

We did manage to construct a GAN model that produced sentences with some resemblance to the distribution it was trying to estimate. Especially, it seemed to be able to end sentences properly (with punctuation, **end-of-sequence-tokens** and padding) and its sentence-construct different from total randomness.

However, the GAN model (when trained with Adam as its optimization algorithm) always seemed to collapse at some point during training, where the generator would no longer be able to get a sentence passed the discriminator. We tried increasing the batch size for the generator and also to stop training the discriminator on encoded random words, but this only alleviated the problem so much. Also, we saw that not using encoded random words preserved the stability of the model much longer, but it became significantly worse at producing just seemingly meaningful sentence-constructs (especially in regards to **end-of-sequence-tokens** and padding).

With SGD, the model would not collapse, but instead it tended towards producing very short sentences with a very, very low non-sensical vocabulary.

All in all, a better sentence encoding model and a more stable GAN model with more optimized hyperparameters and inner architecture would be needed, if the model were to produce sentences of any meaning to a human. As such, our model did not do that.

References

- Chintapalli, Karthik (2019). *Generative Adversarial Networks for Text Generation*. URL: <https://becominghuman.ai/generative-adversarial-networks-for-text-generation-part-1-2b886c8cab10> (visited on 08/13/2020).
- Donahue, David and Anna Rumshisky (2018). “Adversarial Text Generation Without Reinforcement Learning”. In: arXiv: 1810.06640 [cs.CL].
- Goodfellow, Ian J., Jean Pouget-Abadie, M. Mirza, B. Xu, David Warde-Farley, Sherjil Ozair, Aaron C. Courville, and Yoshua Bengio (2014). “Generative Adversarial Networks”. In: *ArXiv* abs/1406.2661.
- Haidar, Md. Akmal and Mehdi Rezagholizadeh (2019). “TextKD-GAN: Text Generation using KnowledgeDistillation and Generative Adversarial Networks”. In: arXiv: 1905.01976 [cs.CL].
- Hui, Jonathan (2018). *GAN — Why it is so hard to train Generative Adversarial Networks!* URL: [https://medium.com/@jonathan_hui/gan-why-it-is-so-hard-to-train-generative-advisory-networks-819a86b3750b](https://medium.com/@jonathan_hui/gan-why-it-is-so-hard-to-train-generative-adversarial-networks-819a86b3750b) (visited on 08/13/2020).
- Inkawhich, Nathan (2020). *DCGAN Tutorial*. URL: https://pytorch.org/tutorials/beginner/dcgan_faces_tutorial.html (visited on 08/13/2020).
- Kusner, Matt J. and José Miguel Hernández-Lobato (2016). “GANS for Sequences of Discrete Elements with the Gumbel-softmax Distribution”. In: arXiv: 1611.04051 [stat.ML].
- Robertson, Sean (2020). *NLP From Scratch: Translation with a Sequence to Sequence Network and Attention*. URL: https://pytorch.org/tutorials/intermediate/seq2seq_translation_tutorial.html (visited on 08/13/2020).
- Sequence-to-Sequence Modeling with nn.Transformer and TorchText* (2020). URL: https://pytorch.org/tutorials/beginner/transformer_tutorial.html (visited on 08/13/2020).
- Silveira, Natalia, Timothy Dozat, Marie-Catherine de Marneffe, Samuel Bowman, Miriam Connor, John Bauer, and Christopher D. Manning (2014). “A Gold Standard Dependency Corpus for English”. In: *Proceedings of the Ninth International Conference on Language Resources and Evaluation (LREC-2014)*.
- Vaswani, Ashish, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin (2017). “Attention Is All You Need”. In: arXiv: 1706.03762 [cs.CL].
- Yu, Lantao, Weinan Zhang, Jun Wang, and Yong Yu (2016). “SeqGAN: Sequence Generative Adversarial Nets with Policy Gradient”. In: arXiv: 1609.05473 [cs.LG].